

TRAINING SIMPLE ML ALGORITHMS FOR CLASSIFICATION

• ARTIFICIAL NEURONS:

- IN the early days of ML McCulloch & Pitts described a Neuron as a simple logic gate.
- They argued that neurons receive some input, and based in this input, modify their output.
- Based on this theory, Rosenblatt published the first perceptron
- The perceptron was able to automatically learn weights & biases that lead the transmission of a signal.

FORMAL DEFINITION OF AN ARTIFICIAL NEURON

- The decision function $\sigma(z)$ takes a linear combination of input values, x , & their corresponding weights, w . (z = the net input ($z = w_1x_1 + w_2x_2 + \dots + w_mx_m$))



$$w = \begin{bmatrix} w_1 \\ \vdots \\ w_m \end{bmatrix} \quad x = \begin{bmatrix} x_1 \\ \vdots \\ x_m \end{bmatrix}$$

- If x_i is greater than some threshold (θ), we predict class 1, otherwise 0.

this is a unit step func. $\sigma(z) = \begin{cases} 1 & \text{if } z \geq \theta \\ 0 & \text{else} \end{cases} \dots \rightarrow z - \theta \geq 0$ can be described as our bias

$$\therefore z = w^T x + b \quad (-\theta)$$

- IN this example, the decision function is the unit step function, which leads to a decision boundary that is linear.



THE PERCEPTRON LEARNING RATE:

ROSENBLATT'S RULE IS AS FOLLOWS:

1. INITIALIZE the weights & bias units to zero or small random #s.
2. FOR EACH TRAINING EXAMPLE x^i
 - a. compute $\hat{y}(i)$ predict
 - b. update weights & biases

TRAINING SIMPLE ML ALGORITHMS FOR CLASSIFICATION

We can formally write the update of the bias unit & weights as:

where

$$w_j := w_j + \Delta w_j \quad \Delta w_j = \eta (y^i - \hat{y}^i) x_j^i$$
$$b := b + \Delta b \quad \Delta b = \eta (y^i - \hat{y}^i)$$

η = learning rate (between 0.0 & 1.0)
 y^i = true class label
 \hat{y}^i = predicted class label

→ the bias unit & all weights are updated simultaneously

→ if the predictions are correct $\Delta w_j = 0$ & $\Delta b = 0$

eg.

$$y_i = 1 \quad \hat{y}_i = 1 \quad \therefore \Delta w_j = \eta (1-1) \cdot x_j^i, \Delta b = \eta (1-1)$$

→ if the predictions are incorrect:


$$y_i = 1 \quad \hat{y}_i = 0 \quad \therefore \Delta w_j = \eta (1-0) \cdot x_j^i = \eta \cdot x_j^i \quad \Delta b = \eta (1-0) = \eta$$

• with binary labels $(y^i - \hat{y}^i) \in \mathbb{R}[-1, 0, 1]$

★ CONVERGENCE OF A PERCEPTRON CAN ONLY HAPPEN IF THE CLASSES ARE LINEARLY SEPERABLE

→ if the classes cant be seperated, we can set a maximum # of epochs (passes over the training dataset) or a threshold for the # of tolerated misclassification.



general concept of a perceptron:



★ SEE THE PERCEPTRON NOTEBOOK

TRAINING SIMPLE ML ALGORITHMS FOR CLASSIFICATION

- The next step more complex is an **ADaptive Linear NEuron (ADALINE)**, where the weights/biases are updated based on a linear activation function, as opposed to a unit step function.
- In adaline the linear activation function is simply the identity of the function $\sigma(z) = z$
- the activation function is used to learn weights, but a threshold function is still used to make a final prediction.



MINIMIZING LOSS FUNCTION w/ GRADIENT DESCENT

- Objective function: loss or cost function we want to minimize
- In Adaline, the loss function (L) is the **Mean Squared Error** between predicted & true outcomes.


$$L(w, b) = \text{MSE} = \frac{1}{n} \sum_{i=1}^n (y^i - \sigma(z^i))^2$$

- The advantage of a linear activation function is that the function can be differentiated.

TRAINING SIMPLE ML ALGORITHMS FOR CLASSIFICATION

GRADIENT DESCENT CONT...

- gradient descent is good @ finding local minimas in our loss function, which will help us determine optimal w & b s.



- using gradient descent, we update the model parameters by taking a step in the opposite direction of the gradient, $\nabla L(w, b)$ of our loss function $L(w, b)$:

$$w := w + \Delta w$$

$$\Delta w = -\eta \nabla_w L(w, b)$$

where:

$$b := b + \Delta b$$

$$\Delta b = -\eta \nabla_b L(w, b)$$

- to compute the gradient of the loss function, we compute the partial derivative w/ respect to each weight

$$\frac{\partial L}{\partial w_j} = -\frac{2}{n} \sum_i (y^i - \sigma(z^i)) \cdot x_j^i$$

PDES!

$$\frac{\partial L}{\partial w_j} = \frac{\partial}{\partial w_j} \cdot \frac{1}{n} \sum_i (y^i - \sigma(z^i))^2$$

$$\frac{\partial L}{\partial b} = -\frac{2}{n} \sum_i (y^i - \sigma(z^i))$$

$$= \frac{1}{n} \frac{\partial}{\partial w_j} \sum_i (y^i - \sigma(z^i))^2$$

$$\Delta w_j = -\eta \cdot \frac{\partial L}{\partial w_j} \quad \& \quad \Delta b = -\eta \frac{\partial L}{\partial b}$$

$$= \frac{2}{n} \sum_i (y^i - \sigma(z^i)) \cdot \frac{\partial}{\partial w_j} (y^i - \sigma(z^i))$$


$$w := w + \Delta w$$

$$b := b + \Delta b$$

$$= \frac{2}{n} \sum_i (y^i - \sigma(z^i)) \left(\frac{\partial y^i}{\partial w_j} - \frac{\partial \sigma(z^i)}{\partial w_j} \right)$$

- THIS IS CLASSIFIED AS BATCH GRADIENT DESCENT BECAUSE THE WEIGHT UPDATE IS BASED ON ALL EXAMPLES IN THE DATASET

if the learning rate is too large:
we will move away from the minima.



TRAINING SIMPLE ML ALGORITHMS FOR CLASSIFICATION

FEATURE SCALING can help our gradient descent loss functions converge.

- **STANDARDIZATION** is a form of feature scaling in which the mean of each feature so that it is centered @ zero & each feature has a standard dev. of 1 (unit variance)

$$x_j' = \frac{x_j - \mu_j}{\sigma_j}$$

$\mu_j \leftarrow \text{mean}$
 $\sigma_j \leftarrow \text{standard deviation}$

- Helpful in gradient descent because weights to converge one unscaled feature on a vastly different scale than other features might destabilize other features.

✳ see Adaline jupyter notebook.

STOCHASTIC GRADIENT DESCENT:

- Normal gradient descent is calculated from the whole training set, however, if our dataset has million of points, this can be very computationally expensive.
- an alternative approach is **stochastic gradient descent (SGD)** where we update the weights & biases incrementally for each training example.

$$\Delta w_j := \eta (y^i - \sigma(z^i)) x_j^i \quad \Delta b = \eta (y^i - \sigma(z^i))$$

- It converges faster than GD, but is slightly more noisy which can help find global minima as opposed to local minima.

✳ IMPORTANT TO SHUFFLE DATA BEFORE SGD

- **ONLINE LEARNING** is where our model is trained on the fly, such as where we are continuously adding data points. Also allows us to discard data after training.